

Unsupervised distributed edge computing person re-identification

Abstract

In this paper, we propose a novel unsupervised adaptive person re-identification method for distributed edge computing environments. The core of the method is to solve the problem of non-independent and identical training data distribution of edge devices in a distributed environment and mining cross-camera pedestrian relationships in a distributed environment. In the field of traditional person re-identification, whenever a new data distribution is encountered, it is necessary to aggregate the data of the labeled source domain and the unlabeled target domain for incremental training. Under the circumstance that the privacy protection of personal data is being paid more attention, it is obviously inappropriate to collect and store pedestrian data on a large scale. Combined with the popular use of edge computing devices such as jetson tx2 in recent years, we set up the following distributed scenarios: edge devices have the ability to train models, edge devices can only get pedestrian data collected by local cameras, and between edge devices No exchange of pedestrian raw image data is allowed. In the above scenario, our method is mainly divided into 4 stages: (1) edge device data style transfer (2) random jump training (3) cross-camera relationship mining and (4) random jump incremental training. Our experiments on the main pedestrian re-identification dataset have proved the effectiveness of the method. In addition, we also analyzed the model training time, inference time, and data interaction bandwidth on the jetson tx2 development board. Prove the practicality of the method

1 Introduction

....

1.1 subsection

....

2 Related work

....

2.1 subsection

....

..

3 Method

Definition In a distributed environment, we adopt the sparse space-time tracklets sampling (SSTT) [Li *et al.*, 2018] to sample the training tracklets from each camera. Each tracklet is represented as $\{t_n^m, y_n^m, g_n^m\}$, where t_n^m means the n -th tracklet in m -th camera. Denoting $t_n^m = \{I_1^{t_n^m}, I_2^{t_n^m}, \dots, I_k^{t_n^m}\}$,

where $I_k^{t_n^m}$ is the k -th image of the n -th tracklet in m -th camera. Each tracklet has two labels, one is a local label and the other is a global label. The local label is represented by y_n^m , whenever the edge device collects new tracklet data, it will automatically be assigned a local label, local labels are distinguishable only in local data and will be used for random walk training. The global label is represented by g_n^m , when distributed clustering is completed, each tracklet will obtain a global label, global labels are distinguishable in global data and will be used for incremental training.

3.1 Single Camera style Transfer

In order to minimize the negative impact of non-Independent and identically distributed (non-IID) in the distributed environment on model performance, each edge device will transform the locally collected data to the target domain through the style transfer. When each edge device is initialized, it will obtain the same pedestrian data as the target domain data for style transfer. Fig.3 describes the above process. we adopt Cycle-GAN [Zhu *et al.*, 2017] as the style transfer model. Given two different domains A and B , the goal of Cycle-GAN is to train two mapping networks G and F so that the distribution of $G(A)$ is as close as possible to the distribution of B and the distribution of $F(B)$ is as close as possible to the distribution of A . The overall CycleGAN loss function is expressed as:

$$\begin{aligned} v(G, F, D_A, D_B) = & v_{GAN}(G, D_B, A, B) \\ & + v_{GAN}(F, D_A, B, A) \\ & + \lambda v_{cyc}(G, F) \end{aligned} \quad (1)$$

where $v_{GAN}(G, D_B, A, B)$ and $v_{GAN}(F, D_A, B, A)$ are the adversarial loss. $v_{cyc}(G, F)$ is the cycle consistency loss. In

order to ensure that the color of the image remains stable before and after the style transfer, we adopt the identity mapping loss [Zhu *et al.*, 2017], which is expressed as:

$$v_{identity}(G, F) = E_{x \sim p_x} [\|F(x) - x\|_1] + E_{y \sim p_y} [\|G(y) - y\|_1] \quad (2)$$

3.2 Random Walk Training

...

3.3 Distributed Clustering

...

3.4 Incremental Training

...

3.5 Global Model Unification

...

4 Experiment

4.1 datasets

All experiments are evaluated on three large-scale image person re-identification datasets (Market-1501 [Zheng *et al.*, 2015], DukeMTMC-reID [Ristani *et al.*, 2016; Zheng *et al.*, 2017] and MSMT17 [Wei *et al.*, 2018]) and one video person re-identification dataset (Mars [Zheng *et al.*, 2016]). Performance is evaluated by the cumulative matching characteristic (CMC) and mean Average Precision (mAP). We adopt the experiments settings and tracklet sampling methods of TAUDL [Li *et al.*, 2018] and UTAL [Li *et al.*, 2019]. For image datasets, we assume all images of a person in one camera are belong to a single tracklet. For video datasets, the same person has multiple tracklets under a single camera, we randomly select one for training.

4.2 Experimental Setup

Style Transfer Model

As introduced in section 3.2, we employ Cycle-GAN [Zhu *et al.*, 2017] as the style transfer model. In each dataset, we choose camera 1 as the target domain for other camera performing a style transfer. we follow the parameter settings of, we resize all images to 256×256 and set the batch size = 1. Beside, the learning rate is 0.0001 for the Discriminator and 0.0002 for the Generator.

Person Re-identification Model

We use ResNet-50 [He *et al.*, 2016] pre-trained on ImageNet [Deng *et al.*, 2009] as our backbone network. Considering the limited computing resources on the edge device, we fix the first two residual layers of ResNet-50, the input image resizes to 256×128 . When training, we use random flip, random cut, random erase [Zhong *et al.*, 2017] for data augmentation. we set the learning rate = 0.1 and use SGD optimizer to train the model. The batch size is set to 0.1. Each edge device needs to be trained for 6 epochs, 8 epochs, 3 epochs, 6 epochs when trained on Market-1501, DukeMTMC-reID, MSMT17 and MARS. In testing, we extract the L2-normalized output of Pooling-5 layer as the image feature and adopt the Euclidean distance to measure the similarities between query and gallery images, which reference [Zhong *et al.*, 2019]

Baseline

We use ResNet50 pre-trained on ImageNet as the experimental baseline.

4.3 Parameter Analysis

We first perform a sensitivity analysis on the hyper-parameter in our experiments. As mentioned in 3.4, the indicator of whether the two memories are of the same type in the clustering process is that they form at least α triangles in the relationship network. We will experiment with different α values in different datasets to detect clustering effects.

Fig 3.1 shows the clustering results when the value of α is 0, 1, 2, 3 on the three datasets of market, duke, and mars. The results show that when the value of α is larger, it is difficult for most memories to establish connected edges with other memories in the relationship network and that means we will get more categories. When α is set to 0, the number of categories after clustering is less than the actual number of categories on all three datasets. Figure 3.2 shows the performance of the model after incremental training using the corresponding clustering results. The performance is optimal when $\alpha = 1$ on the market dataset, the best performance is $\alpha = 2$ on the duke dataset, and on the mars dataset, the performance is optimal when $\alpha = 1$. On three datasets, when $\alpha = 0$, the effect is worse than before the incremental training, which proves that the clustering has no effect.

4.4 Evaluation

Baseline Performance

Table 1 shows the results of the baseline. We tested the model pre-trained on ImageNet directly (called Direct Transfer) on the target testing dataset and found that the effect was very poor. For example, the rank-1 on the market dataset was only 10.0%, and the mAP was only 3.0%. when we train the baseline model on the labeled training set and then test it on the testing set, the baseline (called Supervised Learning) achieves high accuracy.

Ablation experiment

To prove the effectiveness of our proposed method, we conduct ablation studies in Table 1. We first show the effectiveness of random jump training..... Secondly we will show the effectiveness of style transfer..... Finally We will show the effectiveness of clustering.....

Distributed Performance

Run-time analysis We performed model run-time analysis on NVIDIA JETSON TX2. As shown in Table 4, we counted the run-time of the style transfer model inference module, person re-identification model inference module, person re-identification model training module, style transfer model training module and memory retrieval module on jetson tx2. The style transfer model inference module and the person re-identification model inference module count the time it takes to process an image, and the style transfer model training module and the person re-identification training module count the time it takes to train each mini-batch of data.

Communication packet analysis We have statistics on the data packets that the system needs to interact with in a distributed environment. The main types and sizes of data packets are statistically summarized. The statistical results are shown in Table 5.

5 Conclusion

....

6 Acknowledgements

....

References

- [Deng *et al.*, 2009] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale hierarchical image database. In *2009 IEEE conference on computer vision and pattern recognition*, pages 248–255. Ieee, 2009.
- [He *et al.*, 2016] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778, 2016.
- [Li *et al.*, 2018] Minxian Li, Xiatian Zhu, and Shaogang Gong. Unsupervised person re-identification by deep learning tracklet association. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pages 737–753, 2018.
- [Li *et al.*, 2019] Minxian Li, Xiatian Zhu, and Shaogang Gong. Unsupervised tracklet person re-identification. *IEEE transactions on pattern analysis and machine intelligence*, 2019.
- [Ristani *et al.*, 2016] Ergys Ristani, Francesco Solera, Roger Zou, Rita Cucchiara, and Carlo Tomasi. Performance measures and a data set for multi-target, multi-camera tracking. In *European Conference on Computer Vision*, pages 17–35. Springer, 2016.
- [Wei *et al.*, 2018] Longhui Wei, Shiliang Zhang, Wen Gao, and Qi Tian. Person transfer gan to bridge domain gap for person re-identification. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 79–88, 2018.
- [Zheng *et al.*, 2015] Liang Zheng, Liyue Shen, Lu Tian, Shengjin Wang, Jingdong Wang, and Qi Tian. Scalable person re-identification: A benchmark. In *Proceedings of the IEEE international conference on computer vision*, pages 1116–1124, 2015.
- [Zheng *et al.*, 2016] Liang Zheng, Zhi Bie, Yifan Sun, Jingdong Wang, Chi Su, Shengjin Wang, and Qi Tian. Mars: A video benchmark for large-scale person re-identification. In *European Conference on Computer Vision*, pages 868–884. Springer, 2016.
- [Zheng *et al.*, 2017] Zhedong Zheng, Liang Zheng, and Yi Yang. Unlabeled samples generated by gan improve the person re-identification baseline in vitro. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 3754–3762, 2017.
- [Zhong *et al.*, 2017] Zhun Zhong, Liang Zheng, Guoliang Kang, Shaozi Li, and Yi Yang. Random erasing data augmentation. *arXiv preprint arXiv:1708.04896*, 2017.
- [Zhong *et al.*, 2019] Zhun Zhong, Liang Zheng, Zhiming Luo, Shaozi Li, and Yi Yang. Invariance matters: Exemplar memory for domain adaptive person re-identification. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 598–607, 2019.
- [Zhu *et al.*, 2017] Jun-Yan Zhu, Taesung Park, Phillip Isola, and Alexei A Efros. Unpaired image-to-image translation using cycle-consistent adversarial networks. In *Proceedings of the IEEE international conference on computer vision*, pages 2223–2232, 2017.